

C++26 Reflection for JSON Serialization

A Practical Journey

- Daniel Lemire, *University of Quebec* 🇨🇦
- Francisco Geiman Thiesen, *Microsoft* 🇺🇸

CppCon 2025

JSON

- Portable, simple
- Douglas Crockford (2001)
- RFC 8259 (December 2017)

JSON

- scalar values
 - strings (controls and quotes must be escaped)
 - numbers (but not `NaN` or `Inf`)
 - `true` , `false` , `null`
- composed values
 - objects (key/value)
 - arrays (list)

```
{  
  "username": "Alice",  
  "level": 42,  
  "health": 99.5,  
  "inventory": ["sword", "shield", "potion"]  
}
```

JSON downside?

JSON can be *slow*. E.g., 20 MB/s.

- Much slower than disk or network

Micron shows off world's fastest PCIe 6.0 SSD, hitting 27 GB/s speeds — Astera Labs PCIe 6.0 switch enables impressive sequential reads

News

By [Sunny Grimm](#) published March 8, 2025

The next-gen of networking and storage is hitting the trade shows

Performance

- simdjson was the first library to break the gigabyte per second barrier
 - Parsing Gigabytes of JSON per Second, VLDB Journal 28 (6), 2019
 - On-Demand JSON: A Better Way to Parse Documents? SPE 54 (6), 2024
- JSON for Modern C++ can be $100\times$ slower!



Usage

The simdjson library is found in...

- Node.js
- ClickHouse
- Velox



Automate the serialization/deserialization process.



The Problem

Imagine you're building a game server that needs to persist player data.



Player

username: string

level: int

health: double

inventory: string, string, ...

You start simple:

```
struct Player {  
    std::string username;  
    int level;  
    double health;  
    std::vector<std::string> inventory;  
};
```

The Traditional Approach: Manual Serialization

Without reflection, you may write this tedious code:

```
// Serialization - converting Player to JSON
fmt::format(
    "{{"
    "\"username\": \"{}\", "
    "\"level\": {}, "
    "\"health\": {}, "
    "\"inventory\": {}"
    "}}",
    escape_json(p.username),
    p.level,
    std::isfinite(p.health) ? p.health : -1.0,
    p.inventory| std::views::transform(escape_json)
);
```

With a library (JSON for Modern C++)

Or you might use a library.

```
std::string to_json(Player& p) {  
    return nlohmann::json>{"username", p.username},  
                                   {"level", p.level},  
                                   {"health", p.health},  
                                   {"inventory", p.inventory}}  
        .dump();  
}
```

Manual Deserialization (simdjson)

```
object obj = val.get_object();
p.username = obj["username"].get_string();
p.level = obj["username"].get_int64();
p.health = obj["health"].get_double();
array arr = obj["inventory"].get_array();
for (auto item : arr) {
    p.inventory.emplace_back(item.get_string());
}
```

The Pain Points

This manual approach has several problems:

1. **Repetition:** Every field needs to be handled twice (serialize + deserialize)
2. **Maintenance Nightmare:** Add a new field? Update both functions!
3. **Error-Prone:** Typos in field names, forgotten fields, type mismatches
4. **Boilerplate Explosion:** 30+ lines for a simple 4-field struct
5. **Performance:** You may fall into performance traps

When Your Game Grows...

```
struct Equipment {  
    std::string name;  
    int damage; int durability;  
};  
struct Achievement {  
    std::string title; std::string description; bool unlocked;  
    std::chrono::system_clock::time_point unlock_time;  
};  
struct Player {  
    std::string username;  
    int level; double health;  
    std::vector<std::string> inventory;  
    std::map<std::string, Equipment> equipped;           // New!  
    std::vector<Achievement> achievements;               // New!  
    std::optional<std::string> guild_name;               // New!  
};
```

Suddenly you need to write hundreds of lines of serialization code! 🤖

The Solution: C++26 Static Reflection

With C++26 reflection and simdjson, **all that boilerplate disappears:**

```
// Just define your struct - no extra code needed!  
struct Player {  
    std::string username;  
    int level;  
    double health;  
    std::vector<std::string> inventory;  
    std::map<std::string, Equipment> equipped;  
    std::vector<Achievement> achievements;  
    std::optional<std::string> guild_name;  
};
```

Automatic Serialization

```
// Serialization - one line!  
void save_player(const Player& p) {  
    std::string json = simdjson::to_json(p); // That's it!  
    // Save json to file...  
}
```

Automatic Deserialization

```
// Deserialization - one line!  
Player load_player(const std::string& json_str) {  
    return simdjson::from<Player>(json_str); // That's it!  
}
```

- **No manual field mapping**
- **No maintenance burden**
- **Handles nested structures automatically**
- **Performance tuned by the library**

Python

```
# Python
import json
json_str = json.dumps(player.__dict__)
player = Player(**json.loads(json_str))
```



Python reflection

```
def inspect_object(obj):  
    print(f"Class name: {obj.__class__.__name__}")  
    for attr, value in vars(obj).items():  
        print(f"  {attr}: {value}")
```

Go

```
jsonData, err := json.MarshalIndent(player, "", "  ")
if err != nil {
    log.Fatalf("Error during serialization: %v", err)
}
var deserializedPlayer Player
err = json.Unmarshal([]byte(jsonStr), &deserializedPlayer)
```



Go reflection

- Runtime reflection only

```
typ := reflect.TypeOf(obj)
for i := 0; i < typ.NumField(); i++ {
    field := typ.Field(i)
}
```


Java and C#

```
string jsonString = JsonSerializer.Serialize(player, options);  
Player deserializedPlayer = JsonSerializer.Deserialize<Player>(jsonInput, options);
```



Java and C# reflection

- Runtime reflection only.

```
Class<?> playerClass = Player.class;  
Object playerInstance = playerClass.getDeclaredConstructor().newInstance();  
Field nameField = playerClass.getDeclaredField("name");
```

Rust (serde)











```
// Rust with serde  
let json_str = serde_json::to_string(&player)?;  
let player: Player = serde_json::from_str(&json_str)?;
```



Rust reflection

- Rust does not have ANY introspection.
- You cannot enumerate the methods of a struct. Either at runtime or at compile-time.
- Rust relies on annotation (serde) followed by re-parsing of the code.

Reflection as accessing the attributes of a struct.

language	runtime reflection	compile-time reflection
C++ 26		
Go		
Java		
C#		
Rust		

With C++26: simple, maintainable, performant code

```
std::string json_str = simdjson::to_json(player);  
Player player = simdjson::from<Player>(json_str);
```

- **AT COMPILE TIME**
- with no extra tooling
- no annotation

How Does It Work?

The Key Insight: Compile-Time Code Generation

"How can compile-time reflection handle runtime JSON data?"

The answer: Reflection operates on **types and structure**, not runtime values.

It generates regular C++ code at compile time that handles your runtime data.

What Happens Behind the Scenes

```
// What you write:
Player p = simdjson::from<Player>(runtime_json_string);

// What reflection generates at COMPILE TIME (conceptually):
Player deserialize_Player(const json& j) {
    Player p;
    p.username = j["username"].get<std::string>();
    p.level = j["level"].get<int>();
    p.health = j["health"].get<double>();
    p.inventory = j["inventory"].get<std::vector<std::string>>();
    // ... etc for all members
    return p;
}
```


The Actual Reflection Magic

```
template <typename T>
    requires(std::is_class_v<T>) // For user-defined types
error_code deserialize(auto& json_value, T& out) {
    simdjson::ondemand::object obj;
    SIMDJSON_TRY(json_value.get_object().get(obj));

    // This [:expand:] happens at COMPILE TIME
    // It literally generates code for each member
    [:expand(std::meta::nonstatic_data_members_of(^T)):] >> [&<auto member>()] {
        // These are compile-time constants
        constexpr std::string_view field_name = std::meta::identifier_of(member);
        constexpr auto member_type = std::meta::type_of(member);

        // This generates code for each member
        auto err = obj[field_name].get(out.[:member:]);
        if (err && err != simdjson::NO_SUCH_FIELD) {
            return err;
        }
    };

    return simdjson::SUCCESS;
}
```

The `[:expand:]` Statement

The `[:expand:]` statement is the key:

- It's like a **compile-time for-loop**
- Generates code for each struct member
- By the time your program runs, all reflection has been "expanded" into normal C++ code

This means:

- **Zero runtime overhead**
- **Full optimization opportunities**
- **Type safety at compile time**

Compile-Time vs Runtime: What Happens When

```
struct Player {  
    std::string username;    // ← Compile-time: reflection sees this  
    int level;              // ← Compile-time: reflection sees this  
    double health;          // ← Compile-time: reflection sees this  
};  
  
// COMPILE TIME: Reflection reads Player's structure and generates:  
// - Code to read "username" as string  
// - Code to read "level" as int  
// - Code to read "health" as double  
  
// RUNTIME: The generated code processes actual JSON data  
std::string json = R("{\"username\":\"Alice\",\"level\":42,\"health\":100.0}");  
Player p = simdjson::from<Player>(json);  
// Runtime values flow through compile-time generated code
```

Compile-Time Safety: Catching Errors Before They Run

```
// ❌ COMPILE ERROR: Type mismatch detected
struct BadPlayer {
    int username; // Oops, should be string!
};
// simdjson::from<BadPlayer>(json) won't compile if JSON has string

// ❌ COMPILE ERROR: Non-serializable type
struct InvalidType {
    std::thread t; // Threads can't be serialized!
};
// simdjson::to_json(InvalidType{}) fails at compile time

// ✅ COMPILE SUCCESS: All types are serializable
struct GoodType {
    std::vector<int> numbers;
    std::map<std::string, double> scores;
    std::optional<std::string> nickname;
};
```

Zero Overhead: Why It's Fast

Since reflection happens at compile time, there's no runtime penalty:

1. **No runtime type inspection** - everything is known at compile time
2. **No string comparisons for field names** - they become compile-time constants
3. **Optimal code generation** - the compiler sees the full picture
4. **Inline everything** - generated code can be fully optimized

The generated code is often **faster than hand-written code** because:

- It's consistently optimized
- No human errors or inefficiencies
- Leverages simdjson's SIMD parsing throughout

Performance: The Best Part

You might think "automatic = slow", but with simdjson + reflection:

- **Compile-time code generation:** No runtime overhead from reflection
- **SIMD-accelerated parsing:** simdjson uses CPU vector instructions
- **Zero allocation:** String views and in-place parsing
- **Throughput:** ~2-4 GB/s on modern hardware

The generated code is often *faster* than hand-written code!

Real-World Benefits

Before Reflection (Our Game Server example)

- 1000+ lines of serialization code
- Prone to bugs due to serialization mismatching
- Adding new features can imply making tedious changes to boilerplate serialization code

After Reflection

- **0 lines** of serialization code
- **0 serialization bugs** (if it compiles, it works!)
- New features can be added much faster

The Bigger Picture

This pattern extends beyond games:

- **REST APIs**: Automatic request/response serialization
- **Configuration Files**: Type-safe config loading
- **Message Queues**: Serialize/deserialize messages
- **Databases**: Object-relational mapping
- **RPC Systems**: Automatic protocol generation

With C++26 reflection, C++ finally catches up to languages like Rust (serde), Go (encoding/json), and C# (System.Text.Json) in terms of ease of use, but with **better performance** thanks to simdjson's SIMD optimizations.

Try It Yourself

```
struct Meeting {
    std::string title;
    std::chrono::system_clock::time_point start_time;
    std::vector<std::string> attendees;
    std::optional<std::string> location;
    bool is_recurring;
};

// Automatically serializable/deserializable!
std::string json = simdjson::to_json(Meeting{
    .title = "CppCon Planning",
    .start_time = std::chrono::system_clock::now(),
    .attendees = {"Alice", "Bob", "Charlie"},
    .location = "Denver",
    .is_recurring = true
});

Meeting m = simdjson::from<Meeting>(json);
```

Round-Trip Any Data Structure

```
struct TodoItem {
    std::string task;
    bool completed;
    std::optional<std::string> due_date;
};

struct TodoList {
    std::string owner;
    std::vector<TodoItem> items;
    std::map<std::string, int> tags; // tag -> count
};

// Serialize complex nested structures
TodoList my_todos = { /* ... */ };
std::string json = simdjson::to_json(my_todos);

// Deserialize back - perfect round-trip
TodoList restored = simdjson::from<TodoList>(json);
assert(my_todos == restored); // Works if you define operator==
```

The Entire API Surface

Just two functions. Infinite possibilities.

```
simdjson::to_json(object)    // → JSON string  
simdjson::from<T>(json)      // → T object
```

That's it.

No macros. No code generation. No external tools.

Just simdjson leveraging C++26 reflection.

Supporting Standard Containers





Through concepts and template specialization, we support:

- `std::vector<T>` , `std::array<T, N>`
- `std::map<K, V>` , `std::unordered_map<K, V>`
- `std::optional<T>`
- `std::variant<Types...>`
- And many more...

All work seamlessly with reflection!

Conclusion

C++26 Reflection + simdjson =

-  Zero boilerplate
-  Compile-time safety
-  Blazing fast performance
-  Clean, modern API

Welcome to the future of C++ serialization! 

Questions?

Daniel Lemire and Francisco Geiman Thiesen

GitHub: github.com/simdjson/simdjson

Thank you!