# Reflection-based JSON in C++ at Gigabytes per Second

**CppCon 2025**

- Daniel Lemire, *University of Quebec* 🇨🇦
- Francisco Geiman Thiesen, *Microsoft* 🇺🇸

# Why JSON Matters

```json
{
  "age": 5,
  "name": "Daniel",
  "toys": ["wooden dog", "little car"]
}
```

- **Ubiquitous**: Predominant data format on public Web APIs, config files, data interchange

- **Human-readable**: Intuitive attribute-value pairs and arrays

- **MCP**: Format of choice for the design of the MCP protocol

# The C++ JSON Pain

**Performance Bottlenecks:**

- Complex string handling

- Nested structure parsing

- Manual struct ↔️ JSON mapping overhead

- Number parsing and serialization costs

**Safety Pitfalls:**

- Runtime type mismatches

- Unexpected input (broken Unicode)

- No compile-time guarantees

# Other Languages Have It Easy

**Rust (serde):**

```rust
#[derive(Serialize, Deserialize)]
struct Person {
    age: u32,
    name: String,
}
// That's it! Automatic JSON conversion
```

**Similar ergonomics in:** Python, Go, C#, Java, Zig

**C++:** Still manual work... until now

# C++20 Sets the Stage

**Concepts**: Cleaner, constrained interfaces

```cpp
template<Serializable T>
void to_json(const T& obj);
```

**tag_invoke**: Customization point mechanism

**constexpr**: Compile-time computation

*These lay groundwork but fall short of full automation*

# Enter C++26 Reflection (PR2996)

**What it is**: Introspect C++ types at compile time

- Field names, types, structure layout
- "The compiler knows your data structures"

**Why it's game-changing**: No more manual mapping!

**Current state**: Bloomberg's experimental LLVM fork

- github.com/bloomberg/clang-p2996

6

# Simple Reflection Example

```cpp
struct Person {
    int age;
    std::string name;
};

// Reflection automatically exposes:
// - Field count: 2
// - Field names: "age", "name"
// - Field types: int, std::string
// - Automatic JSON mapping possible!
```

# Our Implementation Strategy

**Core Idea**: Reflection auto-generates JSON mapping code

**Serialization**: `Person` → `{"age": 5, "name": "Daniel"}`

- Walk struct fields via reflection

- Emit JSON strings (compile-time optimized)

**Deserialization**: JSON → `Person`

- Parse tokens with type safety

- Populate struct fields automatically

- Comprehensive error handling

# Implementation Architecture

**Three-layer design:**

1. **Lexer**: Tokenize JSON input

2. **Parser**: Build structured representation

3. **Reflection Binder**: Map to/from C++ structs

**Key insight**: Reflection happens at compile time

- Field names become string literals

- Type safety enforced by compiler

- Zero runtime reflection overhead

# First Prototype: Reality Check

```cpp
// Our first attempt was... not impressive
struct Data { /* ... */ };
auto start = std::chrono::high_resolution_clock::now();
Data result = parse_json<Data>(json_string);
auto end = std::chrono::high_resolution_clock::now();
// Result: Much slower than expected!
```

**Problems**: Naive string handling, reflection overhead

**Lesson**: Measure first, optimize second

# Optimization Journey

**Key Performance Wins:**

1. **Compile-time key preparation**

```cpp
constexpr auto field_names = get_field_names<T>();
// Keys computed at compile time!
```

2. **SIMD string operations**

   - Leveraged existing simdjson SIMD expertise

3. **Memory layout optimization**

   - Better cache locality for reflection data

# The SIMD Advantage

```cpp
// Traditional approach
for (char c : json_string) {
    if (c == '"') { /* handle string */ }
    // Character-by-character processing
}

// SIMD approach (simplified)
__m256i chunk = _mm256_loadu_si256(ptr);
__m256i quotes = _mm256_cmpeq_epi8(chunk, quote_chars);
// Process 32 characters at once!
```

*Building on simdjson's proven SIMD foundations*

# Performance Result: The Leap

**Before optimization**: Hundreds of MB/s
**After optimization**: **Gigabytes per second**

*Exact numbers from our ablation study coming...*

**Note**: The profiling experience was... educational

- Some bottlenecks were surprising

- SIMD made the biggest difference

- Compile-time optimization crucial

# Serialization Benchmarks

## Test Setup

- **Hardware**: [To be filled with actual specs]

- **Test data**: Complex nested JSON structures

- **Competitors**: nlohmann/json, RapidJSON, others

## Results Preview

```
Our Implementation:     X.X GB/s
nlohmann/json:         Y.Y GB/s
RapidJSON:             Z.Z GB/s
```

*It's really fast! (But could be faster)*

# Deserialization: Head-to-Head

## The Competition

- **Rust serde**: The gold standard

- **nlohmann/json**: Popular C++ choice

- **RapidJSON**: Performance-focused C++

- **simdjson**: Our non-reflection baseline

## Our Advantage

Reflection + C++ control + SIMD = Winning combination

*Detailed numbers and charts coming in ablation study*

# System Dependencies Matter

**Performance varies by processor:**

- AVX2 vs AVX-512 differences

- Memory bandwidth impact

- Cache architecture effects

**Lesson**: Always benchmark on target hardware
**Good news**: Fast across different architectures

# Safety: What Did We Gain?

**Compile-time Benefits:**

```cpp
struct Person { int age; std::string name; };

// Type mismatch caught at compile time:
Person p = parse_json<Person>(R"({"age": "not_a_number"})");
// ↑ Compiler error, not runtime surprise!
```

**Runtime Benefits:**

- Graceful error handling for malformed JSON

- Clear diagnostic messages

- Structured error reporting

17

# Safety Trade-offs

**Better than handwritten:**

- No manual field mapping bugs

- Type safety enforced automatically

- Consistent error handling

**Not foolproof:**

- Runtime edge cases still possible

- JSON schema evolution challenges

- Unicode edge cases remain tricky

**Net result**: Significant safety improvement

# Downside #1: Learning Curve

**Reflection is new territory:**

```cpp
template<typename T>
constexpr auto serialize(const T& obj) {
    constexpr auto members = std::reflect::members_of(^T);
    // What does ^T mean? How does this work?
    return [=]<std::size_t... Is>(std::index_sequence<Is...>) {
        return make_json_object(
            get_name(members[Is])...,
            obj.[:members[Is]:]...
        );
    }(std::make_index_sequence<members.size()>{});
}
```

**"Template metaprogramming's wild cousin—powerful but untamed"**

# Downside #2: Tooling Challenges

**Debugging is tricky:**

- Reflection errors can be cryptic

- IDE support still catching up

- Template instantiation debugging nightmares

**Example debugging session:**

```
error: no viable conversion from 'std::reflect::info' to 'const char*'
  in instantiation of 'serialize<Person>' requested here
    note: in instantiation of 'get_field_name<0>' requested here
      note: candidate template ignored: substitution failure
```

*Sound familiar, template veterans?*

# Downside #3: Build Times

**The reality check:**

```
# Without reflection
$ time make
real    0m15.234s

# With reflection
$ time make
real    1m23.891s
```

**Why slower:**

- Reflection generates substantial code

- LLVM fork not optimized for compilation speed

- Template instantiation explosion

# Build Time Mitigation

**Strategies we explored:**

1. **Precompiled headers**

```
// precompiled_reflection.h
#include <reflect>
#include "common_types.h"
```

2. **Selective compilation**

   - Only reflection-enable critical paths

   - Traditional parsing for less critical code

3. **Caching strategies**

   - ccache with reflection-aware hashing

# Downside #4: Error Messages

**The brutal truth:**

```
error: invalid use of incomplete type 'std::reflect::member_info<
   std::reflect::get_public_data_members_t<Person>[0]>'
   in instantiation of function template specialization
   'get_member_name<Person, 0>' requested here
     note: in instantiation of function template specialization
     'serialize_impl<Person>' requested here
       note: while substituting template arguments for class template
```

**Real errors from our development**

**Hope**: Compiler vendors will improve diagnostics

# When Should You Try This?

**Today**: Experimental only

- Bloomberg LLVM fork required

- Limited ecosystem support

- Expect rough edges

**Near future (1-2 years)**: Early adopters

- GCC/Clang support incoming

- Better tooling

- Production experiments

**Future (3+ years)**: Mainstream ready

- Full compiler support

# How Beneficial Is It?

**Developer productivity:**

- Eliminate boilerplate JSON mapping

- Automatic type safety

- Consistent error handling

**Performance gains:**

- Gigabyte-per-second parsing

- Compile-time optimizations

- SIMD acceleration built-in

**Maintenance benefits:**

- Schema changes automatically handled

# How Impactful Will This Be?

**For C++ ecosystem:**

- Closes ergonomics gap with other languages

- Enables rapid prototyping with high performance

- Makes C++ more attractive for data-intensive applications

**For your projects:**

- Faster development cycles

- Better performance than hand-written code

- Safer, more maintainable JSON handling

**Game-changer for data-driven C++**

# Code Examples: Before and After

**Traditional C++ approach:**

```cpp
struct Person {
    int age;
    std::string name;
    std::vector<std::string> toys;
};

// Manual JSON handling (20+ lines of boilerplate)
void to_json(json& j, const Person& p) {
    j = json{{"age", p.age}, {"name", p.name}, {"toys", p.toys}};
}
void from_json(const json& j, Person& p) {
    j.at("age").get_to(p.age);
    j.at("name").get_to(p.name);
    j.at("toys").get_to(p.toys);
}
```

# Code Examples: With Reflection

**Reflection-based approach:**

```cpp
struct Person {
    int age;
    std::string name;
    std::vector<std::string> toys;
};

// That's it! Everything else is automatic:
std::string json_str = R"({"age": 5, "name": "Daniel", "toys": ["car"]})";
Person p = simdjson::parse<Person>(json_str);
std::string output = simdjson::serialize(p);
```

**Zero boilerplate. Maximum performance. Compile-time safety.**

# Advanced Features

## Customization hooks:

```cpp
struct Config {
    int timeout_ms;
    std::string endpoint;

    // Custom field naming
    static constexpr auto json_field_names() {
        return std::make_tuple("timeout", "api_endpoint");
    }
};
```

## Validation integration:

```cpp
struct ValidatedData {
    int count;

    void validate() const {
        if (count < 0) throw std::invalid_argument("count must be non-negative");
```

# Future Roadmap

**Short term (2025):**

- Benchmark suite completion

- Error message improvements

- More comprehensive validation

**Medium term (2026-2027):**

- Mainstream compiler support

- Ecosystem integration (fmt, ranges, etc.)

- Production hardening

**Long term (2028+):**

- Standard library integration

# Related Work & Acknowledgments

**Building on giants:**

- simdjson SIMD foundations

- Bloomberg's reflection implementation

- Rust serde design inspiration

- C++20 concepts and customization points

**Community efforts:**

- P2996 reflection proposal authors

- EWG feedback and guidance

- Early adopters and testers

# Practical Considerations

**Memory usage:**

- Reflection metadata is compile-time only

- Runtime overhead minimal

- SIMD requires aligned allocations

**Thread safety:**

- Parser instances are thread-local

- Reflection data is read-only

- Concurrent parsing fully supported

**Integration:**

- Header-only design (when possible)

# Questions for the Audience

**For library authors:**

- How would you integrate this into existing JSON workflows?

- What customization points are most important?

**For application developers:**

- What's your biggest JSON performance bottleneck?

- How much boilerplate would this eliminate for you?

**For the curious:**

- What other domains could benefit from reflection?

# Live Demo

**What we'll show:**

1. Simple struct → JSON conversion

2. Complex nested data parsing

3. Performance comparison

4. Error handling in action

*[Demo section - actual code execution]*

**Note**: Running on Bloomberg LLVM fork

- Reflection syntax may evolve

- Performance representative of potential

# Takeaways

**Key insights:**

1. Reflection + Performance = Possible in C++

2. SIMD makes dramatic difference

3. Compile-time optimization crucial

4. Safety improvements significant

5. Tooling challenges remain

**Bottom line**: The future of C++ JSON handling looks bright
**Call to action**: Try it on Bloomberg's fork, give feedback!

# Resources

**Code & Documentation:**

- Implementation: github.com/simdjson/simdjson/tree/reflection_based_serialization

- Bloomberg LLVM: github.com/bloomberg/clang-p2996

- Reflection proposal: P2996

**Contact:**

- daniel@lemire.me

- franciscogthiesen@gmail.com

*All source code available under business-friendly license*

# Questions?

**Thank you for your attention!**

*Reflection-based JSON in C++ at Gigabytes per Second*

**Ready for the data-driven age**

# Bonus: Implementation Deep Dive

*For the curious - detailed technical discussion*

**Reflection introspection:**

```cpp
template<typename T>
constexpr auto get_field_info() {
    constexpr auto members = std::reflect::members_of(^T);
    return []<std::size_t... Is>(std::index_sequence<Is...>) {
        return std::make_tuple(
            std::reflect::name_of(members[Is])...
        );
    }(std::make_index_sequence<members.size()>{});
}
```

# Bonus: SIMD Integration Details

**String parsing with AVX2:**

```c
__m256i load_and_validate_utf8(__m256i chunk) {
    // Validate UTF-8 sequences
    __m256i high_nibbles = _mm256_and_si256(
        _mm256_srli_epi32(chunk, 4), _mm256_set1_epi8(0x0F));

    // Character classification for JSON
    __m256i whitespace = _mm256_cmpeq_epi8(chunk, _mm256_set1_epi8(' '));
    __m256i quotes = _mm256_cmpeq_epi8(chunk, _mm256_set1_epi8('"'));

    return _mm256_or_si256(whitespace, quotes);
}
```

*This is where the gigabytes per second come from*